

# Heritage Health Claims Data

October 2, 2020

## 1 Predicting Hospital Admissions using Claims Data

### 1.1 Summary

With the rise of readily available data, health-centers nationwide are actively working to minimize costs while producing optimal health outcomes. The only question remains is as follows: how do we predict costs? Hospitalization time is a key driver in healthcare billing. In the **Heritage Health Data Challenge** via Kaggle, I sought to address this via basic classification model building. In addition to exploratory data analysis and feature engineering, I fit three models. The **random forest** algorithm was the most accurate, yielding an accuracy rate of 86.7%

### 1.2 Introduction

Health informatics can carry significant impact with regards to costs and availability of services. The Heritage Health Competition was a past data competition hosted on Kaggle. Participants use available patient data to predict which patients are more likely to experience readmission.

In this project, I use the past datasets to conduct data cleaning, exploratory data analysis, modeling, and appropriate predictive analysis.

### 1.3 Data Processing

The datasets were released via Kaggle in CSV formats. They contain many instances of incomplete cases and require extensive cleaning. The tables were pulled from a relational database, in which the member id is the primary field linking tables. Therefore, joins are required; the **members** and **target** tables have one-to-one relationships, they can be merged using left and/or inner joins. The **drugs** and **labs** tables have a one-to-many relationship with the member table, as they contain records on a yearly basis.

#### 1.3.1 Selecting Predictors

```
[11]:
```

	MemberID	AgeAtFirstClaim	Sex	ClaimsTruncated	DaysInHospital	Year	\
0	210	30-39	NaN	0.0	0.0	Y1	
1	210	30-39	NaN	0.0	0.0	Y3	
2	3197	0-9	F	0.0	0.0	Y1	
3	3197	0-9	F	0.0	0.0	Y2	

4	3197		0-9	F		0.0		0.0	Y3
---	------	--	-----	---	--	-----	--	-----	----

	DrugCount	LabCount	AMI	APPCHOL	...	RENAL2	RENAL3	RESPR4	ROAMI	\
0	2	0	0.0	0.0	...	0.0	0.0	0.0	0.0	
1	2	0	0.0	0.0	...	0.0	0.0	0.0	0.0	
2	1	0	0.0	0.0	...	0.0	0.0	1.0	0.0	
3	2	0	0.0	0.0	...	0.0	0.0	1.0	0.0	
4	1	0	0.0	0.0	...	0.0	0.0	1.0	0.0	

	SEIZURE	SEPSIS	SKNAUT	STROKE	TRAUMA	UTI
0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0

[5 rows x 53 columns]

### 1.3.2 Outcome Variable

[12]:

	MemberID	Year	DSFS	PrimaryConditionGroup	LengthOfStay
1 day	25210	25210	22532		24541
1- 2 weeks	358	358	322		276
2 days	2767	2767	1944		2445
2- 4 weeks	318	318	293		289
26+ weeks	1	1	1		1
3 days	1014	1014	755		848
4 days	418	418	350		312
4- 8 weeks	431	431	418		414
5 days	155	155	129		109
6 days	62	62	55		28

[14]:

	MemberID	Year	DSFS	PrimaryConditionGroup	LengthOfStay	\
0	4	Y2	0- 1 month	RESPR4	NaN	
1	210	Y1	0- 1 month	GIOBSENT	2 days	
3	210	Y1	0- 1 month	GYNEC1	NaN	
4	210	Y1	1- 2 months	MSC2a3	NaN	
6	210	Y1	3- 4 months	PRGNCY	NaN	

	length_recoded
0	0.0
1	2.0
3	0.0
4	0.0
6	0.0

```
[15]:   MemberID Year  length_recoded
      0      4   Y2              0.0
      1     210   Y1              2.0
      2     210   Y2              0.0
      3     210   Y3              0.0
      4    3197   Y1              0.0
```

### 1.3.3 Feature Engineering

One aspect of this project, which may differ from how other participants approached the challenge, entails my experience as a hospital volunteer, a public health student, and later a research assistant. Based on this, rather than employing forward or backward stepwise model building, I will be deliberately selecting features that have documented impacts on health.

One feature that I will be constructing is an SES categorical variable (`low_SES`), derived from the pay delay field. Pay delays can be the result of financial hardship, as I've learned through first hand experience. Socioeconomic status is a key determinant of health and will therefore be included in model building.

Another feature I will be adding is the count of timepoints within a year (`time_count`) in which a patient has a claim. So if a patient has a claim at 0-1 months and 3-4 months during Year One, this feature would be a value of 2.

```
[17]:   MemberID PayDelay
      0      4      43
      1     210      57
      2     210    162+
      3     210     151
      4     210      22
```

```
[19]: (154212, 2)
```

```
[21]:   MemberID Year  DSFS
      0      4   Y2      1
      1     210   Y1      8
      2     210   Y2      6
      3     210   Y3      4
      4    3197   Y1      5
```

### 1.3.4 Merging Datasets Back Together

```
[23]:   MemberID AgeAtFirstClaim  Sex  ClaimsTruncated  DaysInHospital  Year  \
      0     210             30-39  NaN              0.0              0.0      1
      1     210             30-39  NaN              0.0              0.0      3
      2    3197              0-9    F              0.0              0.0      1
      3    3197              0-9    F              0.0              0.0      2
```

4	3197	0-9	F	0.0	0.0	3
5	3713	40-49	F	0.0	0.0	2
6	3741	70-79	F	0.0	0.0	2
7	3889	NaN	F	0.0	0.0	1
8	4048	50-59	M	0.0	0.0	3
9	4187	50-59	F	0.0	0.0	1

	DrugCount	LabCount	AMI	APPCHOL	...	ROAMI	SEIZURE	SEPSIS	SKNAUT	\
0	2.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	
1	2.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	
2	1.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	
3	2.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	
4	1.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	
5	6.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	
6	3.0	5.0	0.0	0.0	...	0.0	0.0	0.0	1.0	
7	3.0	NaN	0.0	0.0	...	0.0	1.0	0.0	0.0	
8	1.0	NaN	0.0	0.0	...	0.0	0.0	0.0	0.0	
9	NaN	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	

	STROKE	TRAUMA	UTI	low_SES	DSFS	length_recoded
0	0.0	0.0	0.0	1.0	8	2.0
1	0.0	0.0	0.0	1.0	4	0.0
2	0.0	0.0	0.0	0.0	5	0.0
3	0.0	0.0	0.0	0.0	5	0.0
4	0.0	0.0	0.0	0.0	11	0.0
5	0.0	0.0	1.0	0.0	10	0.0
6	0.0	0.0	0.0	0.0	20	0.0
7	1.0	0.0	0.0	0.0	13	3.0
8	0.0	0.0	1.0	0.0	22	1.0
9	0.0	0.0	0.0	0.0	4	0.0

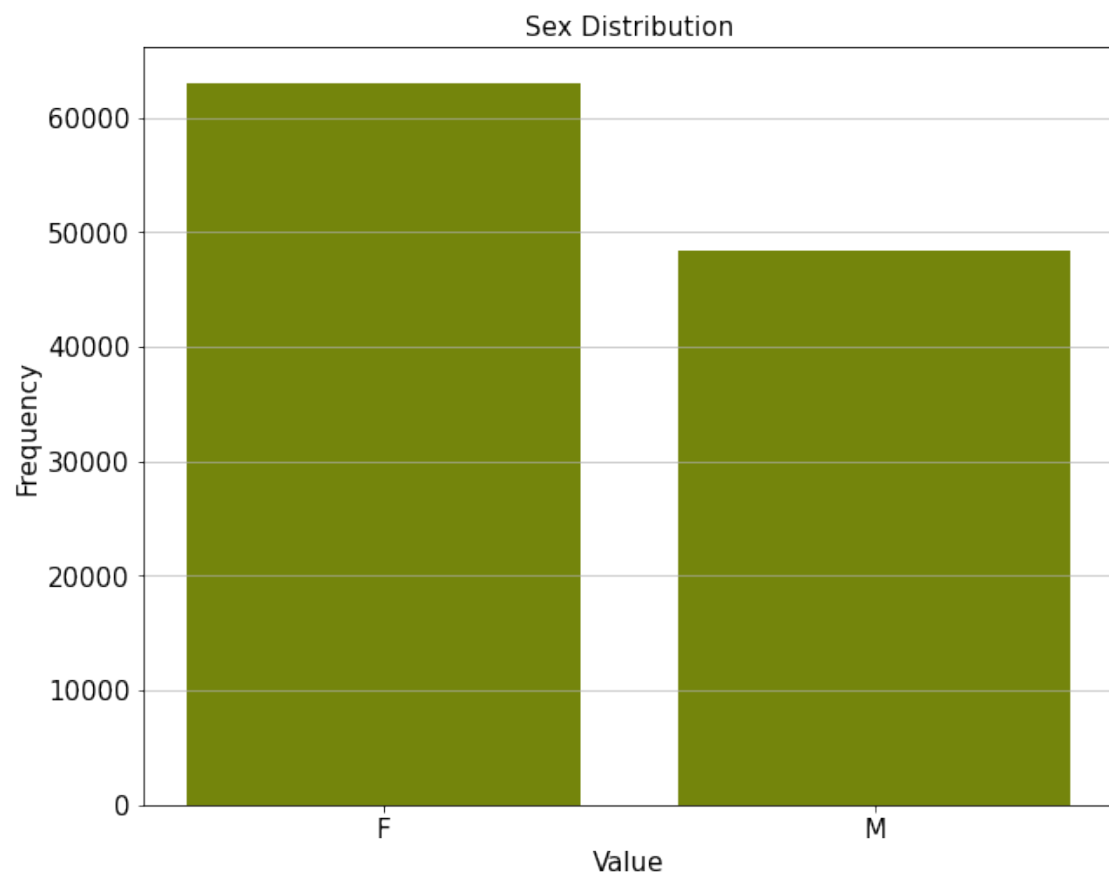
[10 rows x 56 columns]

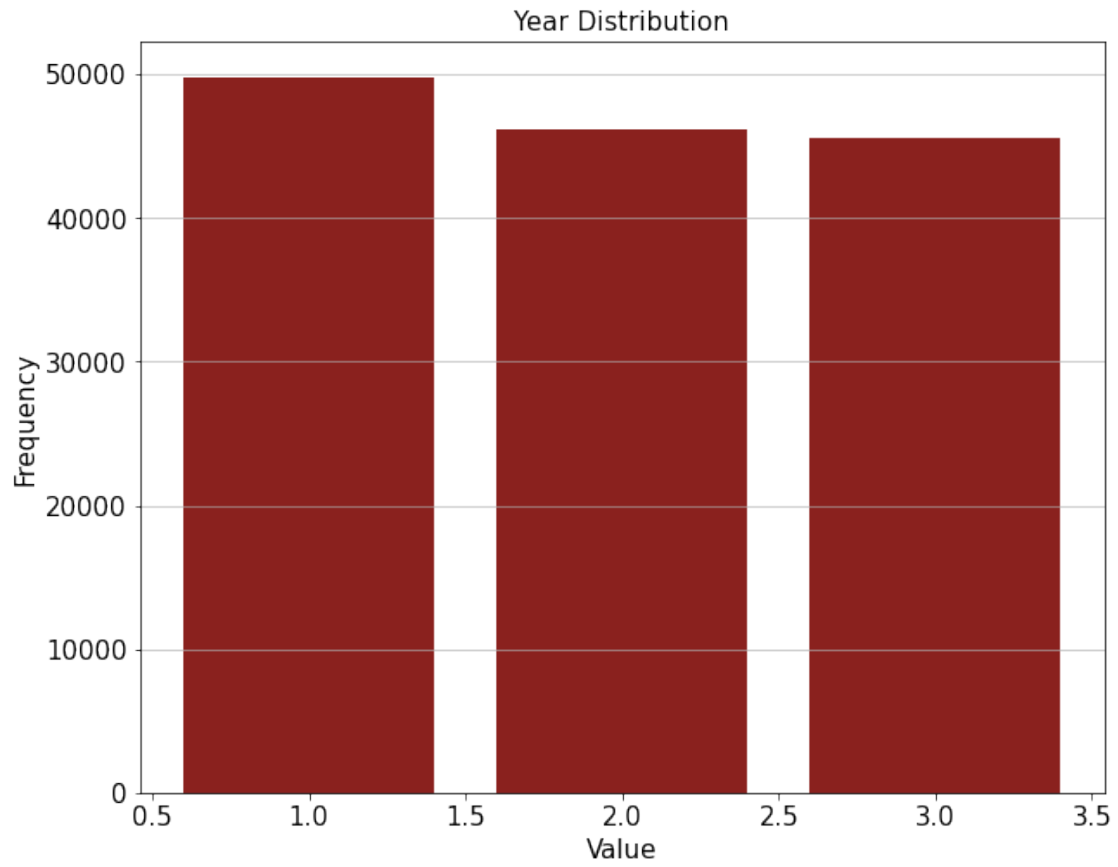
## 1.4 Exploratory Data Analysis

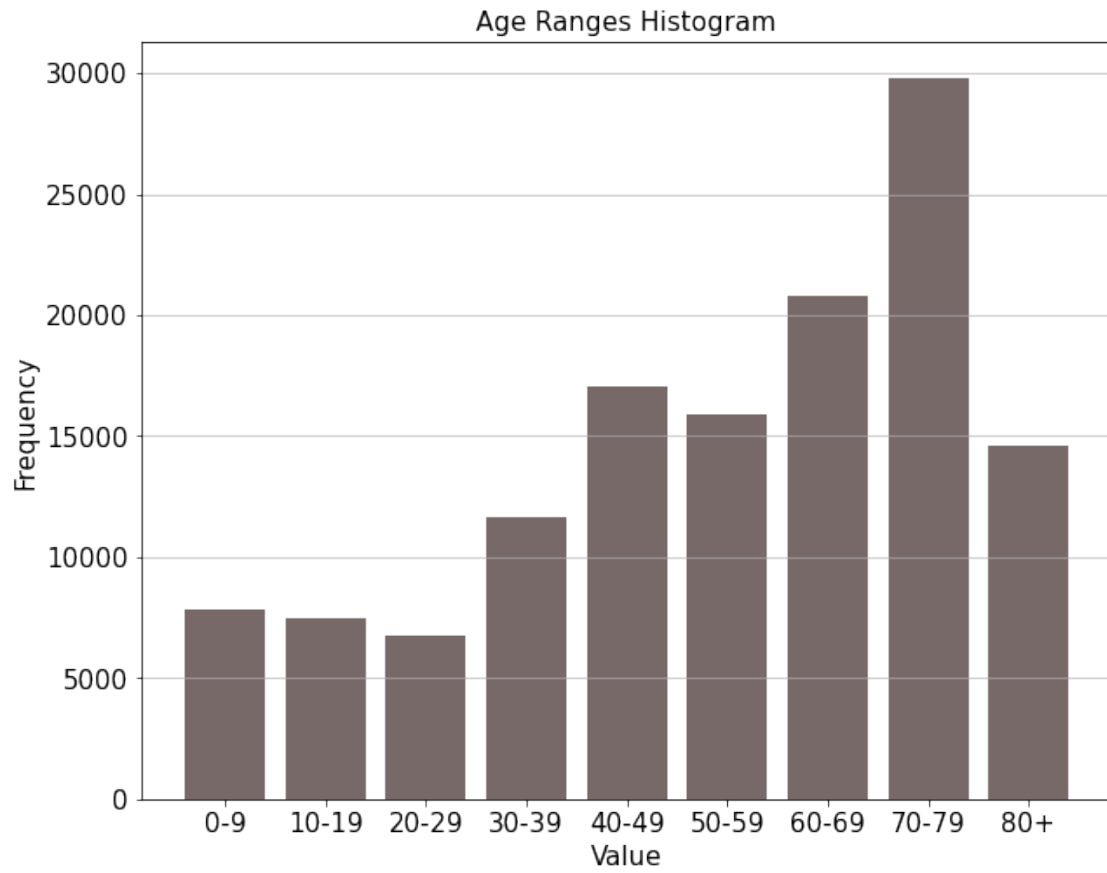
The first step in any data-based problem is understanding the features and outcome we're working with. In addition to visualizing frequencies of specific demographic categories and clinical variables, we'll also visualize the days of hospitalization outcome variable (`length_recoded`).

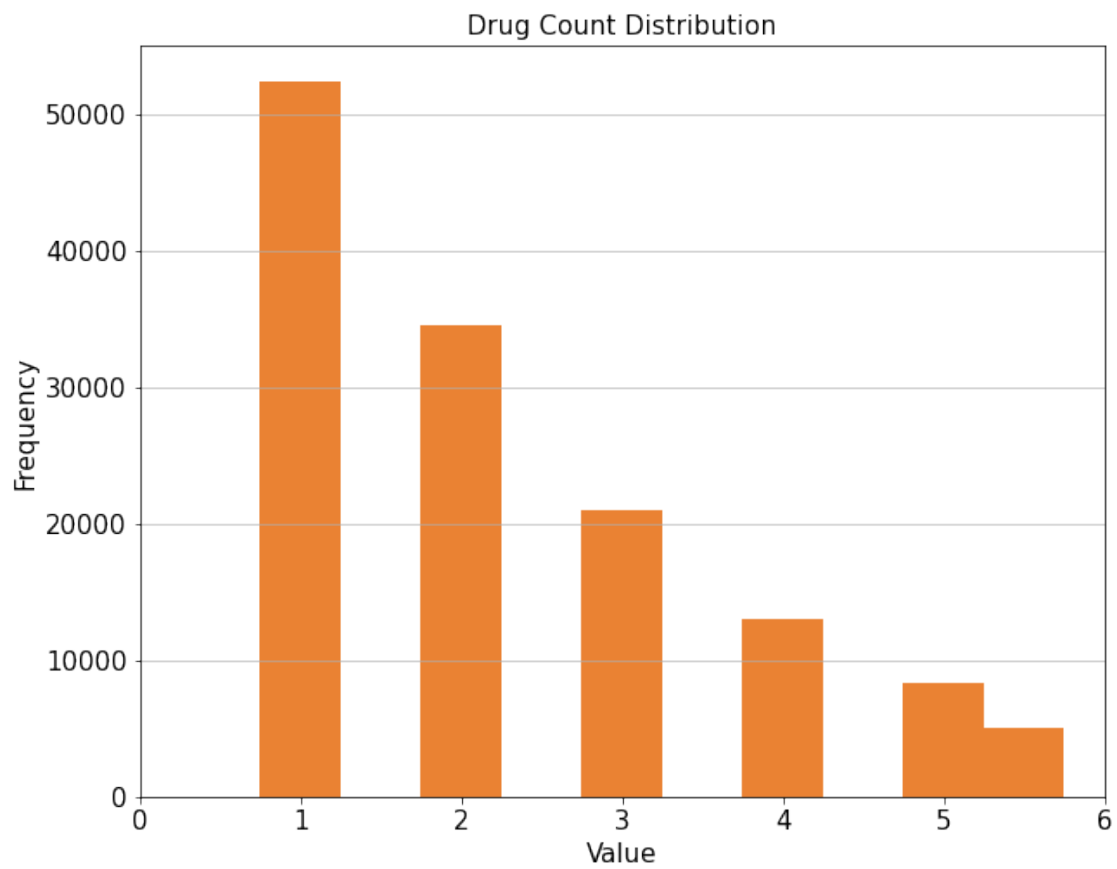
```
[26]: index AgeAtFirstClaim
6      0-9          7848
7     10-19         7478
8     20-29         6756
5     30-39        11663
2     40-49        17041
3     50-59        15862
```

1	60-69	20782
0	70-79	29820
4	80+	14595

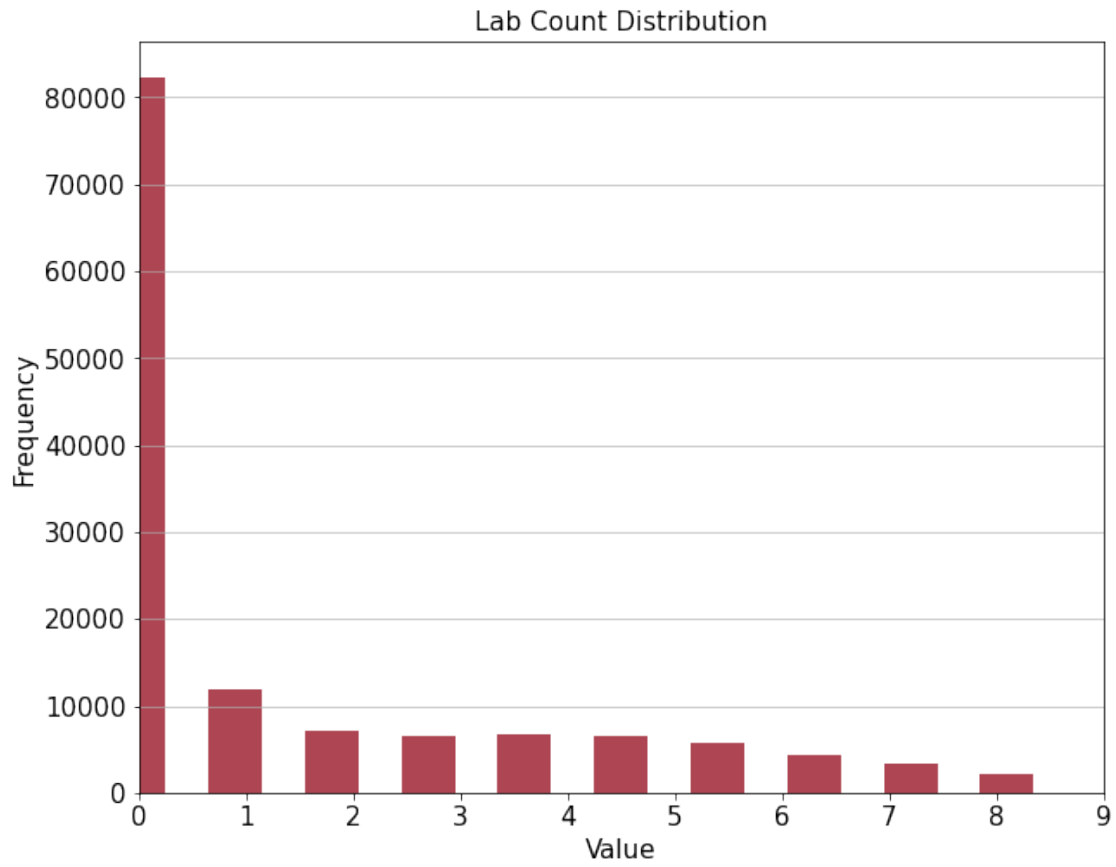






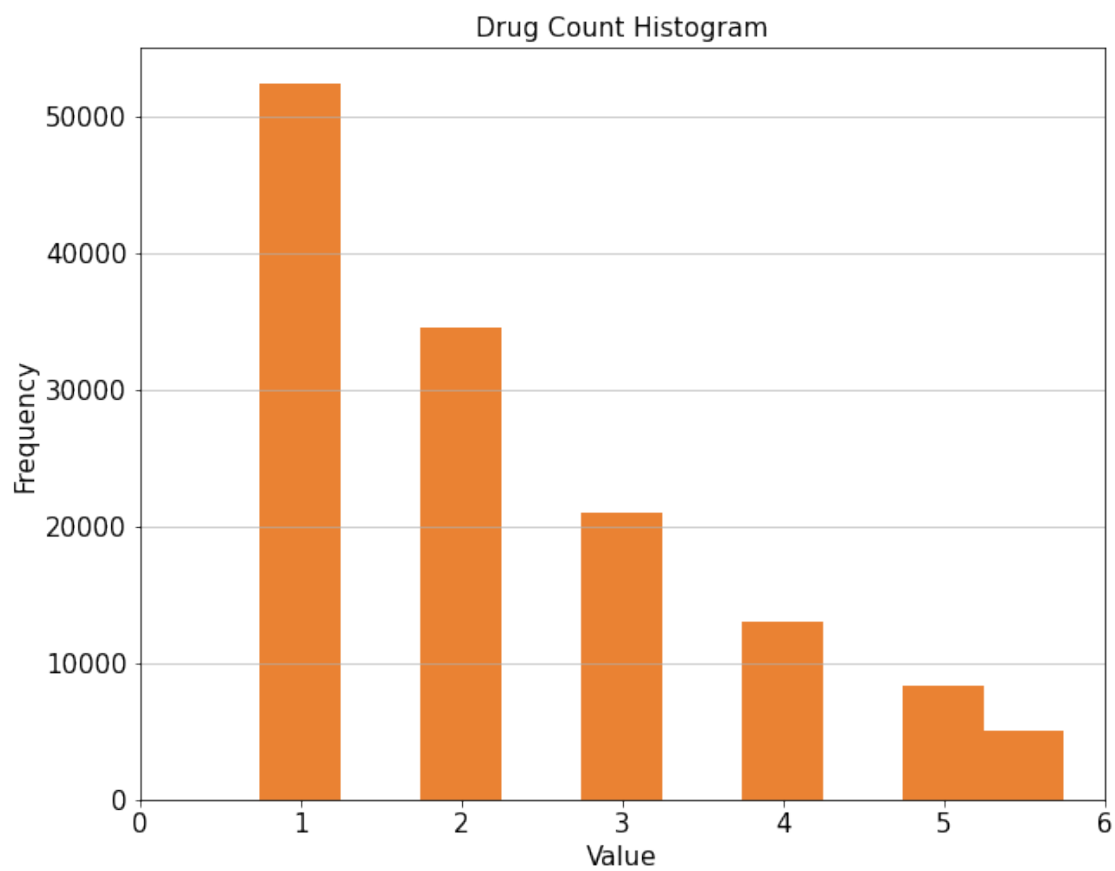







#### 1.4.1 Outcome Visualized

[34]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7ffcda5db5b0>



## 1.5 Modeling

Since we're looking at multiple outcomes with a series of co-variates, decision tree learning would be most appropriate in this scenario. This is further supported knowing that the winning methods used extended decision tree modeling for their predictive analyses.

```
[128]: APPCHOL  ARTHSPIN  CANCRA  CANCRB  CANCRM  CATAST  CHF  COPD  FLaELEC  \
0      0.0      0.0      0.0      0.0      0.0      0.0  0.0  0.0      0.0
1      0.0      0.0      0.0      0.0      0.0      0.0  0.0  0.0      0.0
2      0.0      0.0      0.0      0.0      0.0      0.0  0.0  0.0      0.0

      FXDISLC  ...  DSFS  Year  AgeAtFirstClaim_10-19  AgeAtFirstClaim_20-29  \
0      0.0  ...    8    1                        0                        0
1      0.0  ...    4    3                        0                        0
2      0.0  ...    5    1                        0                        0

      AgeAtFirstClaim_30-39  AgeAtFirstClaim_40-49  AgeAtFirstClaim_50-59  \
0                        1                        0                        0
1                        1                        0                        0
2                        0                        0                        0

      AgeAtFirstClaim_60-69  AgeAtFirstClaim_70-79  AgeAtFirstClaim_80+
0                        0                        0                        0
1                        0                        0                        0
2                        0                        0                        0

[3 rows x 55 columns]
```

```
[40]: count      141558.000000
      mean         0.342220
      std          4.711622
      min          0.000000
      25%          0.000000
      50%          0.000000
      75%          0.000000
      max          441.000000
      Name: length_recoded, dtype: float64
```

```
[105]: (141558, 55)
```

### 1.5.1 Linear Regression: Low SES

While `low_SES` was an engineered feature, based on the patient's pay delay, I was curious what its impact was on other features. This was a question of my own (external to the challenge). I fit an

OLS Linear Regression Model. The low\_SES coefficient, when controlling for age categories, was statistically significantly associated ( = 2.6505,  $p < 0.0001$ ) with the total number of conditions attributed toward hospitalization (per patient). Nothing definitive can be concluded from this, but it is still an interesting observation altogether.

```

                                OLS Regression Results
=====
Dep. Variable:                  sum    R-squared (uncentered):
0.725
Model:                        OLS    Adj. R-squared (uncentered):
0.725
Method:                      Least Squares    F-statistic:
4.672e+04
Date:                        Fri, 02 Oct 2020    Prob (F-statistic):
0.00
Time:                        21:02:31    Log-Likelihood:
-3.3906e+05
No. Observations:            141558    AIC:
6.781e+05
Df Residuals:                141550    BIC:
6.782e+05
Df Model:                    8
Covariance Type:            nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
-----
AgeAtFirstClaim_20-29      2.5657      0.032     79.290      0.000      2.502
2.629
AgeAtFirstClaim_30-39      2.7627      0.025    111.896      0.000      2.714
2.811
AgeAtFirstClaim_40-49      3.0008      0.021    145.983      0.000      2.960
3.041
AgeAtFirstClaim_50-59      3.3048      0.021    154.550      0.000      3.263
3.347
AgeAtFirstClaim_60-69      3.8142      0.019    199.888      0.000      3.777
3.852
AgeAtFirstClaim_70-79      4.2291      0.017    253.426      0.000      4.196
4.262
AgeAtFirstClaim_80+        4.5627      0.023    197.753      0.000      4.517
4.608
low_SES                    2.6505      0.015    172.416      0.000      2.620
2.681
=====
Omnibus:                    9157.372    Durbin-Watson:      1.433

```

Prob(Omnibus):	0.000	Jarque-Bera (JB):	11675.912
Skew:	0.612	Prob(JB):	0.00
Kurtosis:	3.693	Cond. No.	2.80

=====

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

### 1.5.2 Classification Models: Decision Tree

Note that this tree was fit without any dimensionality reduction. As a result, there's definitely room for pruning and making the model more parsimonious. While the tree is ridiculously large and not as helpful as we'd like, the feature importance is worth noting: besides the time variables, RESPR4 (*acute respiratory infections*), ARTHSPIN (*arthropathies and spine disorders*), NEUMENT (*neurological problems*), and low\_SES were ranked the most important features. Overall, the model was 77.4% accurate.

[50]: DecisionTreeClassifier(random\_state=11)

[53]:

	features	importance
45	DSFS	0.188571
46	Year	0.062379
1	ARTHSPIN	0.037378
44	low_SES	0.037099
36	RESPR4	0.035005
26	NEUMENT	0.034198
22	MISCHRT	0.031744
18	INFEC4	0.030325
25	MSC2a3	0.027085
40	SKNAUT	0.026664

Accuracy: 0.774

### 1.5.3 Classification Models: Random Forest

For this model, I utilized a grid search to optimize parameters based on accuracy and refit accordingly. The model was 86.7% accurate. Furthermore, one of the key advantages of random forest was being able to visualize feature importance. GIBLEED and ROAMI were the leading clinical features.

Accuracy: 0.867

[215]:

	name	score
10	GIBLEED	0.15
45	DSFS	0.13
37	ROAMI	0.12

42	TRAUMA	0.09
27	ODaBNCA	0.07